

AI Engineering Job Market Trends

Alexey Grigorev

Eight scrapes, February - August 2026

Today's Agenda

1. Dataset - 6,964 job descriptions across 8 monthly scrapes
2. A data quality fix that changed the numbers
3. What actually moved since February
4. Role archetypes, skills, and the stacks that predict them
5. What still hasn't happened - AI infra, junior hiring

The Dataset

6,964

job descriptions from 2,499 companies

builtin.com, LA, NY, London, Amsterdam, Berlin, India - 8 scrapes, Feb 4 to Aug 25, 2026

Postings per Scrape

Scrape	Jobs
Feb 4	895
Feb 27	869
Mar 27	680
Apr 22	645
May 29	919
Jun 25	888
Jul 22	846
Aug 25	1,222

Zero job-ID overlap between scrapes - each month is an independent cross-section, not a panel

BEFORE THE TRENDS

A Data Quality Fix

[job-market/_internal/eval/README.md](#)

The Extractor Silently Changed

The LLM extraction model switched g_{lm}-5.1 to g_{lm}-5.2 partway through the six original scrapes

Some of what I previously reported as "market trends" was actually the extractor drifting, not the market moving

Re-Extraction, Then a Blind Eval

Rewrote the prompt, re-extracted the full Feb-Jul corpus, then graded old vs. new head-to-head on 50 jobs - blind, random order

Metric	Old	New
Preferred (of 50)	2	38
Classification clean	78.0%	98.0%
Invented skills	74	4

All 8 months now use one consistent, validated extraction

Two "Trends" That Were Actually Noise

- SQL was reported as the fastest-rising skill (+3.59 pp/scrape). On corrected data: a modest +0.5 pp/scrape.
- ai-first share was reported climbing 69.4% → 75.9%. On corrected data: flat, 68-71% the whole time.

Everything from here on uses the corrected, consistent extraction

PART 1

Since February

[job-market/trends.md](#)


ai-first Share Held Steady

71.5% → 70.2%

Feb 4 to Aug 25 - essentially unchanged

The market has held a tight 68-71% ai-first band the whole eight months. No consolidation trend either direction.

The integrator-vs-trainer shift within ai-first roles



Within ai-first jobs, all 8 scrapes

Integrators Are Still Eating Trainers

This is the strongest directional signal in the whole dataset - it got stronger, not weaker, once the extraction was fixed


Stack	Feb 4	Aug 25
Integrator (RAG/agents/APIs)	80.0%	88.0%
Trainer (PyTorch/TF/FT/CUDA)	37.5%	30.4%
RAG + agents together	33.1%	50.5%

PART 2

Role Archetypes

[job-market/trends-appendix.md](#)

Six role archetypes and their skill signatures (k-means, k=6)




Spherical k-means, k=6, deterministic (seed=7)

Six Archetypes - The Split

Archetype	%
RAG app builder	25.3%
Cloud / ML platform engineer	18.0%
Prompt-driven agent builder	14.9%
Pure agent builder	13.6%
DevOps / full-stack engineer	13.3%
ML trainer / researcher	8.2%

Agent-building split into two clusters since the last analysis; DevOps and full-stack merged into one

What companies build with AI (use-case themes)



% of postings whose use cases match, all 8 scrapes

What Companies Build


- Automation / workflow: 52.2% → 59.6% - still the dominant use case
- Agents / autonomous: 43.0% → 48.9%
- RAG / Q&A: 13.6% → 20.6% - up more than half
- Recommend / personalize: 12.8% → 7.0% - fading

PART 3


Skills

Skill trajectories across the six scrapes


Rising skills



Declining skills



Skill momentum: risers (green) vs decliners (red)



Prompt Engineering and RAG Lead Now

Skill	Feb 4	Aug 25	pp/scrape
Prompt Engineering	26.1%	39.0%	+1.46
RAG	32.8%	43.2%	+1.02
LangGraph	7.4%	17.2%	+0.90
MCP	9.9%	17.6%	+0.86

Not SQL - that earlier read was the extractor-switch artifact

The Training Stack Keeps Declining

Skill	Feb 4	Aug 25
PyTorch	20.8%	15.3%
Fine-Tuning	17.1%	12.7%
TensorFlow	12.2%	11.5%

Two Distinct Stacks, Still

Skill co-occurrence lift across all 6,964 jobs

- LangChain → LangGraph 46% of the time (x3.3 lift) - strongest pair in the dataset
- RAG → embeddings (x2.2), vector DBs (x2.0), function calling (x1.9)
- PyTorch → MLOps (x2.2), Kubernetes (x1.4)

MCP and Fine-Tuning now sit closer to the application stack than the training one

Title Predicts the Stack

Title	ai-first	Cust-facing
ai engineer	88%	16%
applied / forward deployed	91%	47%
ml engineer / scientist	74%	10%
platform / infra / devops	47%	11%
data engineer	44%	8%

Applied / forward deployed is the only majority customer-facing group, by a wide margin

PART 4

What Still Hasn't Happened

The "AI Infra Engineer" Still Isn't a Role

19.8% → 9.9%

inference/serving share of ai-first jobs - roughly halved

Inference work keeps getting absorbed into the platform/infra archetype instead of forming its own mass role - a stronger null finding than before

Forward Deployed Engineer

187 postings, 2.7%

Feb 2.0% → peaked 4.4% in July → Aug 2.9%

- 89% customer-facing (vs 14% for other ai-first roles)
- 72% agents, only 47% RAG - agent-heavy generalists
- Only 10% PyTorch/TF - less research, more delivery

The Bottom Rung Is Still Closed

Level	Feb 4	Aug 25
Staff+	18.5%	11.0%
Senior	30.8%	30.3%
Junior/Entry	1.5%	1.2%

Junior hiring has stayed pinned near 1% the entire eight months. Now Staff+ is thinning too.

Key Takeaways

- ai-first share is stable (68-71%), not rising - earlier "consolidation" was measurement noise
- Integrator-vs-trainer shift is real and accelerating: 80% → 88% integrator within ai-first
- Prompt Engineering and RAG are the fastest-rising named skills, not SQL
- Agent-building is fragmenting into more specific sub-roles, not consolidating
- AI infra still isn't a standalone role; junior hiring is still closed

AI Engineering Field Guide

github.com/alexeygrigorev/ai-engineering-field-guide

★ Star the repo to keep an eye on updates

Newsletter: [AI Shipping Blog](#)